

Practice: Modeling Sparse Ratings Using Factorization Machines and the R API

The data set PUBLICEVALS contains information about ratings given by students to courses. A rating is a numeric value between 0 (very bad) and 4 (very good). Not all students take all courses, or even rate them. Therefore, the matrix of ratings is very sparse. The data source contains 88,344 rows and three columns.

| Name | Model Role | Measurement Level | Description |
|---------|------------|-------------------|------------------------------------|
| STUDENT | Input | Nominal | ID of student |
| COURSE | Input | Nominal | ID of course |
| RATING | Target | Nominal | Rating given by student for course |

1. From the Jupyter Lab Home directory, select the plus symbol to open a new page and then select R under Notebook to create a new notebook.
2. Load the SWAT and repr packages. Use the options functions to change the size of the graphical output.
3. Connect to the CAS server and create a connection object.
4. Use the `cas.read.sas7bdat` SWAT function to load the `publicevals.sas7bdat` file onto the CAS server and create a table object reference for the data. Name the new CAS table **evals**.
5. Use SWAT functionality to find the dimension and head of the table, and also the average course rating.
6. Load the simple action set and then use the `distinct` action to find the number of distinct values for all three variables of the **evals** table (student, course, rating). Then, use the `freq` action from the simple action set to find the frequency of each rating level.
7. Load the sampling action set and use the `srs` action to partition the data. Use 90% for training and the remaining 10% for validation. Set the `partind` argument to `TRUE` to add the partition indicator to the CAS table.
8. Build the factorization machine by loading the `factmac` action set and using the `factmac` action. The target for this model is rating and the inputs and nominals are both student and course. Set the maximum iterations and number of factors both to 20 and the `learnstep` for regularization to 0.1. Use the `saveState` and `output` arguments to save the model and the scored training data respectively. Find the MSE in the output for a model fit reference.
9. Score the validation data by first loading the `aStore` action set and then using the `score` action. Pass the score action the validation data using the partition indicator variable and use the `rstore` argument to pass the action the saved factorization machine model. Save the scored validation data as a CAS table and keep all three variables (student, course, rating) in the table using the `copyVars` argument.
10. Find the MSE for the validation data and compare it to the training data MSE. Update the CAS table reference using `defCasTable` and then add the error between the rating and predicted rating as a new variable to the table using open source functionality. Then add the squared error as another variable and use the `mean` function to find the MSE.
11. End the CAS session.